

# merge2

June 14, 2023

## 0.1 Variables are Memory References

We can find the memory address that a variable *references*, by using the `id()` function.

The `id()` function returns the memory address of its argument as a base-10 integer.

We can use the function `hex()` to convert the base-10 number to base-16.

```
[1]: my_var = 10
      print('my_var = {0}'.format(my_var))
      print('memory address of my_var (decimal): {0}'.format(id(my_var)))
      print('memory address of my_var (hex): {0}'.format(hex(id(my_var))))
```

```
my_var = 10
memory address of my_var (decimal): 1968827120
memory address of my_var (hex): 0x7559eaf0
```

```
[2]: greeting = 'Hello'
      print('greeting = {0}'.format(greeting))
      print('memory address of my_var (decimal): {0}'.format(id(greeting)))
      print('memory address of my_var (hex): {0}'.format(hex(id(greeting))))
```

```
greeting = Hello
memory address of my_var (decimal): 1688681719264
memory address of my_var (hex): 0x1892d4625e0
```

Note how the memory address of `my_var` is **different** from that of `greeting`.

Strictly speaking, `my_var` is not “equal” to 10.

Instead `my_var` is a **reference** to an (*integer*) object (*containing the value 10*) located at the memory address `id(my_var)`

Similarly for the variable `greeting`.

## 0.2 Reference Counting

Method that returns the reference count for a given variable’s memory address:

```
[1]: import ctypes

      def ref_count(address):
```

```
return ctypes.c_long.from_address(address).value
```

Let's make a variable, and check it's reference count:

```
[2]: my_var = [1, 2, 3, 4]
     ref_count(id(my_var))
```

```
[2]: 1
```

There is another built-in function we can use to obtain the reference count:

```
[3]: import sys
     sys.getrefcount(my_var)
```

```
[3]: 2
```

But why is this returning 2, instead of the expected 1 we obtained with the previous function?

Answer: The `sys.getrefcount()` function takes **my\_var** as an argument, this means it receives (and stores) a reference to **my\_var**'s memory address **also** - hence the count is off by 1. So we will use `from_address()` instead.

We make another reference to the **same** reference as **my\_var**:

```
[4]: other_var = my_var
```

Let's look at the memory address of those two variables and the reference counts:

```
[5]: print(hex(id(my_var)), hex(id(other_var)))
     print(ref_count(id(my_var)))
```

```
0x1e43f368388 0x1e43f368388
2
```

Force one reference to go away:

```
[6]: other_var = None
```

And we look at the reference count again:

```
[7]: print(ref_count(id(my_var)))
```

```
1
```

We see that the reference count has gone back to 1.

You'll probably never need to do anything like this in Python. Memory management is completely transparent - this is just to illustrate some of what is going behind the scenes as it helps to understand upcoming concepts.

### 0.3 Garbage Collection

```
[1]: import ctypes
import gc
```

We use the same function that we used in the lesson on reference counting to calculate the number of references to a specified object (using its memory address to avoid creating an extra reference)

```
[2]: def ref_count(address):
    return ctypes.c_long.from_address(address).value
```

We create a function that will search the objects in the GC for a specified id and tell us if the object was found or not:

```
[3]: def object_by_id(object_id):
    for obj in gc.get_objects():
        if id(obj) == object_id:
            return "Object exists"
    return "Not found"
```

Next we define two classes that we will use to create a circular reference

Class A's constructor will create an instance of class B and pass itself to class B's constructor that will then store that reference in some instance variable.

```
[4]: class A:
    def __init__(self):
        self.b = B(self)
        print('A: self: {0}, b:{1}'.format(hex(id(self)), hex(id(self.b))))
```

```
[5]: class B:
    def __init__(self, a):
        self.a = a
        print('B: self: {0}, a: {1}'.format(hex(id(self)), hex(id(self.a))))
```

We turn off the GC so we can see how reference counts are affected when the GC does not run and when it does (by running it manually).

```
[6]: gc.disable()
```

Now we create an instance of A, which will, in turn, create an instance of B which will store a reference to the calling A instance.

```
[7]: my_var = A()
```

```
B: self: 0x1fc1eae44e0, a: 0x1fc1eae4908
A: self: 0x1fc1eae4908, b:0x1fc1eae44e0
```

As we can see A and B's constructors ran, and we also see from the memory addresses that we have a circular reference.

In fact `my_var` is also a reference to the same `A` instance:

```
[8]: print(hex(id(my_var)))
```

```
0x1fc1eae4908
```

Another way to see this:

```
[9]: print('a: \t{0}'.format(hex(id(my_var))))  
     print('a.b: \t{0}'.format(hex(id(my_var.b))))  
     print('b.a: \t{0}'.format(hex(id(my_var.b.a))))
```

```
a:      0x1fc1eae4908  
a.b:    0x1fc1eae44e0  
b.a:    0x1fc1eae4908
```

```
[10]: a_id = id(my_var)  
      b_id = id(my_var.b)
```

We can see how many references we have for `a` and `b`:

```
[11]: print('refcount(a) = {0}'.format(ref_count(a_id)))  
     print('refcount(b) = {0}'.format(ref_count(b_id)))  
     print('a: {0}'.format(object_by_id(a_id)))  
     print('b: {0}'.format(object_by_id(b_id)))
```

```
refcount(a) = 2  
refcount(b) = 1  
a: Object exists  
b: Object exists
```

As we can see the `A` instance has two references (one from `my_var`, the other from the instance variable `b` in the `B` instance)

The `B` instance has one reference (from the `A` instance variable `a`)

Now, let's remove the reference to the `A` instance that is being held by `my_var`:

```
[12]: my_var= None
```

```
[13]: print('refcount(a) = {0}'.format(ref_count(a_id)))  
     print('refcount(b) = {0}'.format(ref_count(b_id)))  
     print('a: {0}'.format(object_by_id(a_id)))  
     print('b: {0}'.format(object_by_id(b_id)))
```

```
refcount(a) = 1  
refcount(b) = 1  
a: Object exists  
b: Object exists
```

As we can see, the reference counts are now both equal to 1 (a pure circular reference), and reference counting alone did not destroy the A and B instances - they're still around. If no garbage collection is performed this would result in a memory leak.

Let's run the GC manually and re-check whether the objects still exist:

```
[14]: gc.collect()
      print('refcount(a) = {0}'.format(ref_count(a_id)))
      print('refcount(b) = {0}'.format(ref_count(b_id)))
      print('a: {0}'.format(object_by_id(a_id)))
      print('b: {0}'.format(object_by_id(b_id)))
```

```
refcount(a) = 0
refcount(b) = 0
a: Not found
b: Not found
```

### 0.3.1 Dynamic Typing

Python is dynamically typed.

This means that the type of a variable is simply the type of the object the variable name points to (references). The variable itself has no associated type.

```
[1]: a = "hello"
```

```
[2]: type(a)
```

```
[2]: str
```

```
[3]: a = 10
```

```
[4]: type(a)
```

```
[4]: int
```

```
[5]: a = lambda x: x**2
```

```
[6]: a(2)
```

```
[6]: 4
```

```
[7]: type(a)
```

```
[7]: function
```

As you can see from the above examples, the type of the variable `a` changed over time - in fact it was simply the type of the object `a` was referencing at that time. No type was ever attached to the variable name itself.

```
[ ]:
```

## 0.4 Variable Re-Assignment

Notice how the memory address of **a** is different every time.

```
[1]: a = 10
     hex(id(a))
```

```
[1]: '0x7559eaf0'
```

```
[2]: a = 15
     hex(id(a))
```

```
[2]: '0x7559eb90'
```

```
[3]: a = 5
     hex(id(a))
```

```
[3]: '0x7559ea50'
```

```
[4]: a = a + 1
     hex(id(a))
```

```
[4]: '0x7559ea70'
```

However, look at this:

```
[5]: a = 10
     b = 10
     print(hex(id(a)))
     print(hex(id(b)))
```

```
0x7559eaf0
```

```
0x7559eaf0
```

The memory addresses of both **a** and **b** are the same!!

We'll revisit this in a bit to explain what is going on.

## 0.5 Object Mutability

Certain Python built-in object types (aka data types) are **mutable**.

That is, the internal contents (state) of the object in memory can be modified.

```
[1]: my_list = [1, 2, 3]
     print(my_list)
     print(hex(id(my_list)))
```

```
[1, 2, 3]
0x1cf6ab5b208
```

```
[2]: my_list.append(4)
      print(my_list)
      print(hex(id(my_list)))
```

```
[1, 2, 3, 4]
0x1cf6ab5b208
```

As you can see, the memory address of *my\_list* has **not** changed.

But, the **contents** of *my\_list* has changed from *[1, 2, 3]* to *[1, 2, 3, 4]*.

On the other hand, consider this:

```
[3]: my_list_1 = [1, 2, 3]
      print(my_list_1)
      print(hex(id(my_list_1)))
```

```
[1, 2, 3]
0x1cf6abd55c8
```

```
[4]: my_list_1 = my_list_1 + [4]
      print(my_list_1)
      print(hex(id(my_list_1)))
```

```
[1, 2, 3, 4]
0x1cf6ab56888
```

Notice here that the memory address of *my\_list\_1* **did** change.

This is because concatenating two lists objects *my\_list\_1* and *[4]* did not modify the contents of *my\_list\_1* - instead it created a new list object and re-assigned *my\_list\_1* to reference this new object.

Similarly with **dictionary** objects that are also **mutable** types.

```
[5]: my_dict = dict(key1='value 1')
      print(my_dict)
      print(hex(id(my_dict)))
```

```
{'key1': 'value 1'}
0x1cf6abdc8c8
```

```
[6]: my_dict['key1'] = 'modified value 1'
      print(my_dict)
      print(hex(id(my_dict)))
```

```
{'key1': 'modified value 1'}
0x1cf6abdc8c8
```

```
[7]: my_dict['key2'] = 'value 2'
      print(my_dict)
      print(hex(id(my_dict)))
```

```
{'key1': 'modified value 1', 'key2': 'value 2'}
0x1cf6abdc8c8
```

Once again we see that while we are modifying the **contents** of the dictionary, the memory address of *my\_dict* has not changed.

Now consider the immutable sequence type: **tuple**

The tuple is immutable, so elements cannot be added, removed or replaced.

```
[8]: t = (1, 2, 3)
```

This tuple will **never** change at all. It has three elements, the integers 1, 2, and 3. This will remain the case as long as **t**'s reference is not changed.

But, consider the following tuple:

```
[9]: a = [1, 2]
      b = [3, 4]
      t = (a, b)
```

Now, **t** is still immutable, i.e. it contains a reference to the object **a** and the object **b**. **That** will never change as long as **t**'s reference is not re-assigned.

**However**, the elements **a** and **b** are, themselves, mutable.

```
[10]: a.append(3)
       b.append(5)
       print(t)
```

```
([1, 2, 3], [3, 4, 5])
```

Observe that the contents of **a** and **b** **did** change!

So immutability can be a little more subtle than just thinking something can never change.

The tuple **t** did **not** change - it contains two elements, that are the references **a** and **b**. And that will not change. But, because the referenced elements are mutable themselves, it appears as though the tuple has changed.

It hasn't though - that distinction is subtle but important to understand!

## 0.6 Function Arguments and Mutability

Consider a function that receives a *string* argument, and changes the argument in some way:

```
[1]: def process(s):
      print('initial s # = {}'.format(hex(id(s))))
      s = s + ' world'
```



```
print('s after change # = {0}'.format(hex(id(s))))
```

```
[2]: my_var = 'hello'
      print('my_var # = {0}'.format(hex(id(my_var))))
```

my\_var # = 0x1e7e96fc420

Note that when *s* is received, it is referencing the same object as *my\_var*.

After we “modify” *s*, *s* is pointing to a new memory address:

```
[3]: process(my_var)
```

initial s # = 0x1e7e96fc420

s after change # = 0x1e7e97153b0

And our own variable *my\_var* is still pointing to the original memory address:

```
[4]: print('my_var # = {0}'.format(hex(id(my_var))))
```

my\_var # = 0x1e7e96fc420

Let’s see how this works with mutable objects:

```
[5]: def modify_list(items):
      print('initial items # = {0}'.format(hex(id(items))))
      if len(items) > 0:
          items[0] = items[0] ** 2
      items.pop()
      items.append(5)
      print('final items # = {0}'.format(hex(id(items))))
```

```
[6]: my_list = [2, 3, 4]
      print('my_list # = {0}'.format(hex(id(my_list))))
```

my\_list # = 0x1e7e972d308

```
[7]: modify_list(my_list)
```

initial items # = 0x1e7e972d308

final items # = 0x1e7e972d308

```
[8]: print(my_list)
      print('my_list # = {0}'.format(hex(id(my_list))))
```

[4, 3, 5]

my\_list # = 0x1e7e972d308

As you can see, throughout all the code, the memory address referenced by *my\_list* and *items* is always the **same** (shared) reference - we are simply modifying the contents (**internal state**) of the object at that memory address.

Now, even with immutable container objects we have to be careful, e.g. a tuple containing a list (the tuple is immutable, but the list element inside the tuple **is** mutable)

```
[9]: def modify_tuple(t):  
      print('initial t # = {0}'.format(hex(id(t))))  
      t[0].append(100)  
      print('final t # = {0}'.format(hex(id(t))))
```

```
[10]: my_tuple = ([1, 2], 'a')
```

```
[11]: hex(id(my_tuple))
```

```
[11]: '0x1e7e9614288'
```

```
[12]: modify_tuple(my_tuple)
```

```
initial t # = 0x1e7e9614288  
final t # = 0x1e7e9614288
```

```
[13]: my_tuple
```

```
[13]: ([1, 2, 100], 'a')
```

As you can see, the first element of the tuple was mutated.

## 0.7 Shared References and Mutability

The following sets up a shared reference between the variables `my_var_1` and `my_var_2`

```
[1]: my_var_1 = 'hello'  
      my_var_2 = my_var_1  
      print(my_var_1)  
      print(my_var_2)
```

```
hello  
hello
```

```
[2]: print(hex(id(my_var_1)))  
      print(hex(id(my_var_2)))
```

```
0x24c9144ca08  
0x24c9144ca08
```

```
[3]: my_var_2 = my_var_2 + ' world!'
```

```
[4]: print(hex(id(my_var_1)))  
      print(hex(id(my_var_2)))
```

```
0x24c9144ca08
0x24c9144fab0
```

Be careful if the variable type is mutable!

Here we create a list (*my\_list\_1*) and create a variable (*my\_list\_2*) referencing the same list object:

```
[5]: my_list_1 = [1, 2, 3]
      my_list_2 = my_list_1
      print(my_list_1)
      print(my_list_2)
```

```
[1, 2, 3]
[1, 2, 3]
```

As we can see they have the same memory address (shared reference):

```
[6]: print(hex(id(my_list_1)))
      print(hex(id(my_list_2)))
```

```
0x24c9144fc48
0x24c9144fc48
```

Now we modify the list referenced by *my\_list\_2*:

```
[7]: my_list_2.append(4)
```

*my\_list\_2* has been modified:

```
[8]: print(my_list_2)
```

```
[1, 2, 3, 4]
```

And since *my\_list\_1* references the same list object, it has also changed:

```
[9]: print(my_list_1)
```

```
[1, 2, 3, 4]
```

As you can see, both variables still share the same reference:

```
[10]: print(hex(id(my_list_1)))
       print(hex(id(my_list_2)))
```

```
0x24c9144fc48
0x24c9144fc48
```

## 0.8 ### Behind the scenes with Python's memory manager

Recall from a few lectures back:

```
[11]: a = 10
      b = 10
```

```
[12]: print(hex(id(a)))  
      print(hex(id(b)))
```

```
0x7559eaf0  
0x7559eaf0
```

Same memory address!!

This is safe for Python to do because integer objects are **immutable**.

So, even though *a* and *b* initially shared the same memory address, we can never modify *a*'s value by “modifying” *b*'s value.

The only way to change *b*'s value is to change its reference, which will never affect *a*.

```
[13]: b = 15
```

```
[14]: print(hex(id(a)))  
      print(hex(id(b)))
```

```
0x7559eaf0  
0x7559eb90
```

However, for mutable objects, Python's memory manager does not do this, since that would **not** be safe.

```
[15]: my_list_1 = [1, 2, 3]  
      my_list_2 = [1, 2, 3]
```

As you can see, although the two variables were assigned identical “contents”, the memory addresses are not the same:

```
[16]: print(hex(id(my_list_1)))  
      print(hex(id(my_list_2)))
```

```
0x24c9146c5c8  
0x24c913c6848
```

## 0.9 Variable Equality

From the previous lecture we know that **a** and **b** will have a **shared** reference:

```
[1]: a = 10  
     b = 10  
  
     print(hex(id(a)))  
     print(hex(id(b)))
```

```
0x7559eaf0  
0x7559eaf0
```

When we use the **is** operator, we are comparing the memory address **references**:

```
[2]: print("a is b: ", a is b)
```

a is b: True

But if we use the == operator, we are comparing the **contents**:

```
[3]: print("a == b:", a == b)
```

a == b: True

The following however, do not have a shared reference:

```
[4]: a = [1, 2, 3]
     b = [1, 2, 3]

     print(hex(id(a)))
     print(hex(id(b)))
```

0x27006f17288

0x27006e968c8

Although they are not the same objects, they do contain the same “values”:

```
[5]: print("a is b: ", a is b)
     print("a == b", a == b)
```

a is b: False

a == b True

Python will attempt to compare values as best as possible, for example:

```
[6]: a = 10
     b = 10.0
```

These are **not** the same reference, since one object is an **int** and the other is a **float**

```
[7]: print(type(a))
     print(type(b))
```

<class 'int'>

<class 'float'>

```
[8]: print(hex(id(a)))
     print(hex(id(b)))
```

0x7559eaf0

0x270064b1870

```
[9]: print('a is b:', a is b)
     print('a == b:', a == b)
```

```
a is b: False
a == b: True
```

So, even though *a* is an integer 10, and *b* is a float 10.0, the values will still compare as equal.

In fact, this will also have the same behavior:

```
[10]: c = 10 + 0j
      print(type(c))
```

```
<class 'complex'>
```

```
[11]: print('a is c:', a is c)
      print('a == c:', a == c)
```

```
a is c: False
a == c: True
```

## 0.10 ### The None Object

**None** is a built-in “variable” of type *NoneType*.

Basically the keyword **None** is a reference to an object instance of *NoneType*.

*NoneType* objects are immutable! Python’s memory manager will therefore use shared references to the *None* object.

```
[12]: print(None)
```

```
None
```

```
[13]: hex(id(None))
```

```
[13]: '0x75576bc0'
```

```
[14]: type(None)
```

```
[14]: NoneType
```

```
[15]: a = None
      print(type(a))
      print(hex(id(a)))
```

```
<class 'NoneType'>
0x75576bc0
```

```
[16]: a is None
```

```
[16]: True
```

```
[17]: a == None
```

```
[17]: True
```

```
[18]: b = None  
      hex(id(b))
```

```
[18]: '0x75576bc0'
```

```
[19]: a is b
```

```
[19]: True
```

```
[20]: a == b
```

```
[20]: True
```

```
[21]: l = []
```

```
[22]: type(l)
```

```
[22]: list
```

```
[23]: l is None
```

```
[23]: False
```

```
[24]: l == None
```

```
[24]: False
```

## 0.11 Everything is an Object

```
[1]: a = 10
```

**a** is an object of type **int**, i.e. **a** is an instance of the **int** class.

```
[2]: print(type(a))
```

```
<class 'int'>
```

If **int** is a class, we should be able to declare it using standard class instantiation:

```
[3]: b = int(10)
```

```
[4]: print(b)  
      print(type(b))
```

```
10
```

```
<class 'int'>
```

We can even request the class documentation:

```
[5]: help(int)
```

Help on class int in module builtins:

```
class int(object)
|   int(x=0) -> integer
|   int(x, base=10) -> integer
|
|   Convert a number or string to an integer, or return 0 if no arguments
|   are given.  If x is a number, return x.__int__().  For floating point
|   numbers, this truncates towards zero.
|
|   If x is not a number or if base is given, then x must be a string,
|   bytes, or bytearray instance representing an integer literal in the
|   given base.  The literal can be preceded by '+' or '-' and be surrounded
|   by whitespace.  The base defaults to 10.  Valid bases are 0 and 2-36.
|   Base 0 means to interpret the base from the string as an integer literal.
|   >>> int('0b100', base=0)
|   4
|
|   Methods defined here:
|
|   __abs__(self, /)
|       abs(self)
|
|   __add__(self, value, /)
|       Return self+value.
|
|   __and__(self, value, /)
|       Return self&value.
|
|   __bool__(self, /)
|       self != 0
|
|   __ceil__(...)
|       Ceiling of an Integral returns itself.
|
|   __divmod__(self, value, /)
|       Return divmod(self, value).
|
|   __eq__(self, value, /)
|       Return self==value.
|
|   __float__(self, /)
|       float(self)
```



```

|  __floor__(...)
|      Flooring an Integral returns itself.
|
|  __floordiv__(self, value, /)
|      Return self//value.
|
|  __format__(...)
|      default object formatter
|
|  __ge__(self, value, /)
|      Return self>=value.
|
|  __getattr__(self, name, /)
|      Return getattr(self, name).
|
|  __getnewargs__(...)
|
|  __gt__(self, value, /)
|      Return self>value.
|
|  __hash__(self, /)
|      Return hash(self).
|
|  __index__(self, /)
|      Return self converted to an integer, if self is suitable for use as an
index into a list.
|
|  __int__(self, /)
|      int(self)
|
|  __invert__(self, /)
|      ~self
|
|  __le__(self, value, /)
|      Return self<=value.
|
|  __lshift__(self, value, /)
|      Return self<<value.
|
|  __lt__(self, value, /)
|      Return self<value.
|
|  __mod__(self, value, /)
|      Return self%value.
|
|  __mul__(self, value, /)
|      Return self*value.
|

```

```

|  __ne__(self, value, /)
|      Return self!=value.
|
|  __neg__(self, /)
|      -self
|
|  __new__(*args, **kwargs) from builtins.type
|      Create and return a new object.  See help(type) for accurate signature.
|
|  __or__(self, value, /)
|      Return self|value.
|
|  __pos__(self, /)
|      +self
|
|  __pow__(self, value, mod=None, /)
|      Return pow(self, value, mod).
|
|  __radd__(self, value, /)
|      Return value+self.
|
|  __rand__(self, value, /)
|      Return value&self.
|
|  __rdivmod__(self, value, /)
|      Return divmod(value, self).
|
|  __repr__(self, /)
|      Return repr(self).
|
|  __rfloordiv__(self, value, /)
|      Return value//self.
|
|  __rlshift__(self, value, /)
|      Return value<<self.
|
|  __rmod__(self, value, /)
|      Return value%self.
|
|  __rmul__(self, value, /)
|      Return value*self.
|
|  __ror__(self, value, /)
|      Return value|self.
|
|  __round__(...)
|      Rounding an Integral returns itself.
|      Rounding with an ndigits argument also returns an integer.

```

```

|
| __rpow__(self, value, mod=None, /)
|     Return pow(value, self, mod).
|
| __rrshift__(self, value, /)
|     Return value>>self.
|
| __rshift__(self, value, /)
|     Return self>>value.
|
| __rsub__(self, value, /)
|     Return value-self.
|
| __rtruediv__(self, value, /)
|     Return value/self.
|
| __rxor__(self, value, /)
|     Return value^self.
|
| __sizeof__(...)
|     Returns size in memory, in bytes
|
| __str__(self, /)
|     Return str(self).
|
| __sub__(self, value, /)
|     Return self-value.
|
| __truediv__(self, value, /)
|     Return self/value.
|
| __trunc__(...)
|     Truncating an Integral returns itself.
|
| __xor__(self, value, /)
|     Return self^value.
|
| bit_length(...)
|     int.bit_length() -> int
|
|     Number of bits necessary to represent self in binary.
|     >>> bin(37)
|     '0b100101'
|     >>> (37).bit_length()
|     6
|
| conjugate(...)
|     Returns self, the complex conjugate of any int.

```

```

|
| from_bytes(...) from builtins.type
|     int.from_bytes(bytes, byteorder, *, signed=False) -> int
|
|     Return the integer represented by the given array of bytes.
|
|     The bytes argument must be a bytes-like object (e.g. bytes or
bytearray).
|
|     The byteorder argument determines the byte order used to represent the
|     integer. If byteorder is 'big', the most significant byte is at the
|     beginning of the byte array. If byteorder is 'little', the most
|     significant byte is at the end of the byte array. To request the native
|     byte order of the host system, use `sys.byteorder' as the byte order
value.
|
|     The signed keyword-only argument indicates whether two's complement is
|     used to represent the integer.
|
| to_bytes(...)
|     int.to_bytes(length, byteorder, *, signed=False) -> bytes
|
|     Return an array of bytes representing an integer.
|
|     The integer is represented using length bytes. An OverflowError is
|     raised if the integer is not representable with the given number of
|     bytes.
|
|     The byteorder argument determines the byte order used to represent the
|     integer. If byteorder is 'big', the most significant byte is at the
|     beginning of the byte array. If byteorder is 'little', the most
|     significant byte is at the end of the byte array. To request the native
|     byte order of the host system, use `sys.byteorder' as the byte order
value.
|
|     The signed keyword-only argument determines whether two's complement is
|     used to represent the integer. If signed is False and a negative
integer
|     is given, an OverflowError is raised.
|
| -----
| Data descriptors defined here:
|
| denominator
|     the denominator of a rational number in lowest terms
|
| imag
|     the imaginary part of a complex number

```

```
|
| numerator
|     the numerator of a rational number in lowest terms
|
| real
|     the real part of a complex number
```

As we see from the docs, we can even create an **int** using an overloaded constructor:

```
[6]: b = int('10', base=2)
```

```
[7]: print(b)
      print(type(b))
```

```
2
<class 'int'>
```

## 0.12 ### Functions are Objects too

```
[8]: def square(a):
      return a ** 2
```

```
[9]: type(square)
```

```
[9]: function
```

In fact, we can even assign them to a variable:

```
[10]: f = square
```

```
[11]: type(f)
```

```
[11]: function
```

```
[12]: f is square
```

```
[12]: True
```

```
[13]: f(2)
```

```
[13]: 4
```

```
[14]: type(f(2))
```

```
[14]: int
```

A function can return a function

```
[15]: def cube(a):  
      return a ** 3
```

```
[16]: def select_function(fn_id):  
      if fn_id == 1:  
          return square  
      else:  
          return cube
```

```
[17]: f = select_function(1)  
      print(hex(id(f)))  
      print(hex(id(square)))  
      print(hex(id(cube)))  
      print(type(f))  
      print('f is square: ', f is square)  
      print('f is cube: ', f is cube)  
      print(f)  
      print(f(2))
```

```
0x21257457b70  
0x21257457b70  
0x21255fab8c8  
<class 'function'>  
f is square: True  
f is cube: False  
<function square at 0x0000021257457B70>  
4
```

```
[18]: f = select_function(2)  
      print(hex(id(f)))  
      print(hex(id(square)))  
      print(hex(id(cube)))  
      print(type(f))  
      print('f is square: ', f is square)  
      print('f is cube: ', f is cube)  
      print(f)  
      print(f(2))
```

```
0x21255fab8c8  
0x21257457b70  
0x21255fab8c8  
<class 'function'>  
f is square: False  
f is cube: True  
<function cube at 0x0000021255FAB8C8>  
8
```

We could even call it this way:

```
[19]: select_function(1)(5)
```

```
[19]: 25
```

A Function can be passed as an argument to another function  
(This example is pretty useless, but it illustrates the point effectively)

```
[20]: def exec_function(fn, n):  
      return fn(n)
```

```
[21]: result = exec_function(cube, 2)  
      print(result)
```

```
8
```

We will come back to functions as arguments **many** more times throughout this course!

### 0.13 Python Optimizations: Interning

Earlier, we saw shared references being created automatically by Python:

```
[1]: a = 10  
     b = 10  
     print(id(a))  
     print(id(b))
```

```
1968827120  
1968827120
```

Note how `a` and `b` reference the same object.

But consider the following example:

```
[2]: a = 500  
     b = 500  
     print(id(a))  
     print(id(b))
```

```
1935322088624  
1935322089008
```

As you can see, the variables `a` and `b` do **not** point to the same object!

This is because Python pre-caches integer objects in the range `[-5, 256]`

So for example:

```
[3]: a = 256  
     b = 256  
     print(id(a))  
     print(id(b))
```

```
1968834992
1968834992
```

and

```
[4]: a = -5
      b = -5
      print(id(a))
      print(id(b))
```

```
1968826640
1968826640
```

do have the same reference.

This is called **interning**: Python **interns** the integers in the range [-5, 256].

The integers in the range [-5, 256] are essentially **singleton** objects.

```
[5]: a = 10
      b = int(10)
      c = int('10')
      d = int('1010', 2)
```

```
[6]: print(a, b, c, d)
```

```
10 10 10 10
```

```
[7]: a is b
```

```
[7]: True
```

```
[8]: a is c
```

```
[8]: True
```

```
[9]: a is d
```

```
[9]: True
```

As you can see, all these variables were created in different ways, but since the integer object with value 10 behaves like a singleton, they all ended up pointing to the **same** object in memory.

## 0.14 Python Optimizations: String Interning

Python will automatically intern *certain* strings.

In particular all the identifiers (variable names, function names, class names, etc) are interned (singleton objects created).

Python will also intern string literals that look like identifiers.



For example:

```
[1]: a = 'hello'
      b = 'hello'
      print(id(a))
      print(id(b))
```

1342722069536

1342722069536

But not the following:

```
[2]: a = 'hello, world!'
      b = 'hello, world!'
      print(id(a))
      print(id(b))
```

1342722047024

1342722170928

However, because the following literals resemble identifiers, even though they are quite long, Python will still automatically intern them:

```
[3]: a = 'hello_world'
      b = 'hello_world'
      print(id(a))
      print(id(b))
```

1342722047856

1342722047856

And even longer:

```
[4]: a = '_this_is_a_long_string_that_could_be_used_as_an_identifier'
      b = '_this_is_a_long_string_that_could_be_used_as_an_identifier'
      print(id(a))
      print(id(b))
```

1342721886784

1342721886784

Even if the string starts with a digit:

```
[5]: a = '1_hello_world'
      b = '1_hello_world'
      print(id(a))
      print(id(b))
```

1342722046256

1342722046256

That was interned (pointer is the same), but look at this one:

```
[6]: a = '1 hello world'
     b = '1 hello world'
     print(id(a))
     print(id(b))
```

```
1342722046832
```

```
1342722172592
```

Interning strings (making them singleton objects) means that testing for string equality can be done faster by comparing the memory address:

```
[7]: a = 'this_is_a_long_string'
     b = 'this_is_a_long_string'
     print('a==b:', a == b)
     print('a is b:', a is b)
```

```
a==b: True
```

```
a is b: True
```

**Note: Remember, using `is` ONLY works if the strings were interned!** Here's where this technique fails:

```
[8]: a = 'hello world'
     b = 'hello world'
     print('a==b:', a==b)
     print('a is b:', a is b)
```

```
a==b: True
```

```
a is b: False
```

You *can* force strings to be interned (but only use it if you have a valid performance optimization need):

```
[9]: import sys
```

```
[10]: a = sys.intern('hello world')
      b = sys.intern('hello world')
      c = 'hello world'
      print(id(a))
      print(id(b))
      print(id(c))
```

```
1342722172080
```

```
1342722172080
```

```
1342722174896
```

Notice how `a` and `b` are pointing to the same object, but `c` is **NOT**.

So, since both `a` and `b` were interned we can use `is` to test for equality of the two strings:

```
[11]: print('a==b:', a==b)
      print('a is b:', a is b)
```

```
a==b: True
a is b: True
```

So, does interning really make a big speed difference?

Yes, but only if you are performing a *lot* of comparisons.

Let's run some quick and dirty benchmarks:

```
[12]: def compare_using_equals(n):
      a = 'a long string that is not interned' * 200
      b = 'a long string that is not interned' * 200
      for i in range(n):
          if a == b:
              pass
```

```
[13]: def compare_using_interning(n):
      a = sys.intern('a long string that is not interned' * 200)
      b = sys.intern('a long string that is not interned' * 200)
      for i in range(n):
          if a is b:
              pass
```

```
[14]: import time

      start = time.perf_counter()
      compare_using_equals(10000000)
      end = time.perf_counter()

      print('equality: ', end-start)
```

```
equality:  2.965451618090112
```

```
[15]: start = time.perf_counter()
      compare_using_interning(10000000)
      end = time.perf_counter()

      print('identity: ', end-start)
```

```
identity:  0.28690104431129626
```

As you can see, the performance difference, especially for long strings, and for many comparisons, can be quite radical!

## 0.15 Python Peephole Optimizations

Peephole optimizations refer to a certain class of optimization strategies Python employs during any compilation phases.

**Constant Expressions** Let's see how Python reduces constant expressions for optimization purposes:

```
[20]: def my_func():
      a = 24 * 60
      b = (1, 2) * 5
      c = 'abc' * 3
      d = 'ab' * 11
      e = 'the quick brown fox' * 10
      f = [1, 2] * 5
```

```
[21]: my_func.__code__.co_consts
```

```
[21]: (None,
      24,
      60,
      1,
      2,
      5,
      'abc',
      3,
      'ab',
      11,
      'the quick brown fox',
      10,
      1440,
      (1, 2),
      (1, 2, 1, 2, 1, 2, 1, 2, 1, 2),
      'abcabcabc')
```

As you can see in the example above, `24 * 60` was pre-calculated and cached as a constant (1440). Similarly, `(1, 2) * 5` was cached as `(1, 2, 1, 2, 1, 2, 1, 2, 1, 2)` and `'abc' * 3` was cached as `abcabcabc`.

On the other hand, note how `'the quick brown fox' * 10` was **not** pre-calculated (too long). Similarly `[1, 2] * 5` was not pre-calculated either since a list is *mutable*, and hence not a *constant*.

**Membership Tests** In membership testing, optimizations are applied as can be seen below:

```
[69]: def my_func():
      if e in [1, 2, 3]:
          pass
```

```
[70]: my_func.__code__.co_consts
```

```
[70]: (None, 1, 2, 3, (1, 2, 3))
```

As you can see, the mutable list `[1, 2, 3]` was converted to an immutable tuple.

It is OK to do this here, since we are testing membership of the list **at that point in time**, hence it is safe to convert it to a tuple, which is more efficient than testing membership of a list.

In the same way, set membership will be converted to frozen set membership:

```
[22]: def my_func():  
      if e in {1, 2, 3}:  
          pass
```

```
[23]: my_func.__code__.co_consts
```

```
[23]: (None, 1, 2, 3, frozenset({1, 2, 3}))
```

In general, when you are writing your code, if you can use `set` membership testing, prefer that over a list or tuple - it is quite a bit more efficient.

Let's do a small quick (and dirty) benchmark of this:

```
[5]: import string  
import time  
  
char_list = list(string.ascii_letters)  
char_tuple = tuple(string.ascii_letters)  
char_set = set(string.ascii_letters)  
  
print(char_list)  
print()  
print(char_tuple)  
print()  
print(char_set)
```

```
['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p',  
'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'A', 'B', 'C', 'D', 'E', 'F',  
'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V',  
'W', 'X', 'Y', 'Z']
```

```
('a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p',  
'q', 'r', 's', 't', 'u', 'v', 'w', 'x', 'y', 'z', 'A', 'B', 'C', 'D', 'E', 'F',  
'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V',  
'W', 'X', 'Y', 'Z')
```

```
{'l', 'p', 'x', 'R', 'j', 'S', 's', 'T', 'W', 'Y', 'Z', 'P', 'g', 'O', 'b', 'u',  
'H', 'G', 'v', 'e', 'M', 'n', 'w', 't', 'Q', 'E', 'N', 'X', 'C', 'i', 'A', 'B',
```

```
'F', 'V', 'a', 'm', 'r', 'f', 'h', 'U', 'D', 'c', 'y', 'z', 'J', 'd', 'o', 'I',  
'L', 'K', 'k', 'q']
```

```
[6]: def membership_test(n, container):  
      for i in range(n):  
          if 'p' in container:  
              pass
```

```
[7]: start = time.perf_counter()  
      membership_test(10000000, char_list)  
      end = time.perf_counter()  
      print('list membership: ', end-start)
```

```
list membership: 2.6035404184015434
```

```
[8]: start = time.perf_counter()  
      membership_test(10000000, char_tuple)  
      end = time.perf_counter()  
      print('tuple membership: ', end-start)
```

```
tuple membership: 2.602491734651276
```

```
[9]: start = time.perf_counter()  
      membership_test(10000000, char_set)  
      end = time.perf_counter()  
      print('set membership: ', end-start)
```

```
set membership: 0.3743007599607324
```

As you can see, set membership tests run quite a bit faster - which is not surprising since they are basically dictionary-like objects, so hash maps are used for looking up an item to determine membership.